



Very Large Scale Integration II - VLSI II

Cache & Memory Systems

Prof. Dr. Berna Örs Yalçın

Res. Asst. Serdar Duran

ITU VLSI Laboratories
Istanbul Technical University




What is a Cache?

- All programs spend much of their time for accessing memory.
- Therefore, **memory system** is a **major factor** for the general computer performance.
- In order to improve computer performance, there are some mechanisms including **memory hierarchy** and **cache** design.



What is a Cache?

- In a computer system, the second lowest level of the memory hierarchy is a cache. First is the registers.
- The aim of the cache is that make the access memory faster while keep the memory size large!





Principle of Locality

- **The principle of locality** states that programs access a small portion of the address space repetitively over a short period of time!
- Two different types of locality:
 - **Temporal locality:** the reuse of specific data within a small time duration.
 - **Spatial locality:** the use of data elements within close memory locations.



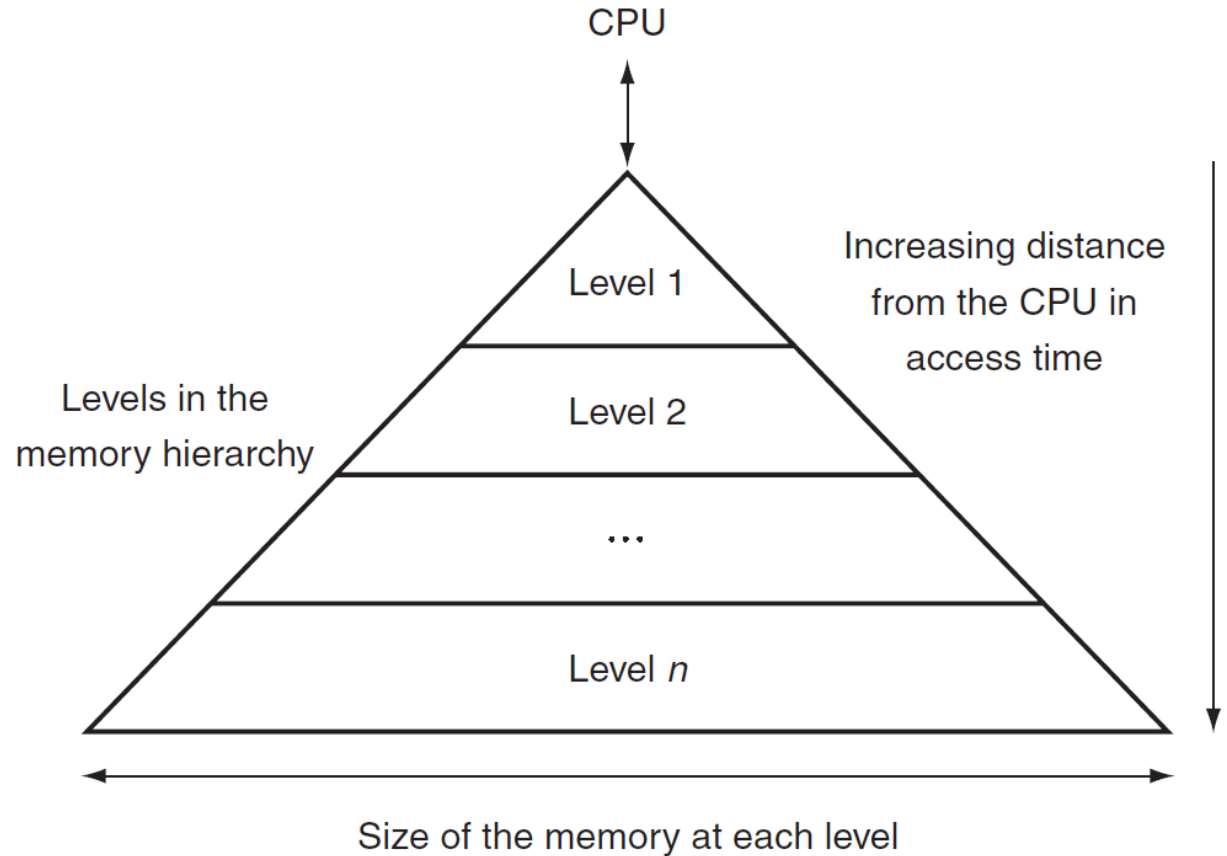
Principle of Locality

- **Temporal Locality:** Most programs contain loops, so instructions and data are likely to be accessed repeatedly, showing high amounts of temporal locality.
- **Spatial Locality:** Since instructions are normally accessed sequentially (no branching), programs show high spatial locality. Besides, accessing to elements of an array or a record have high spatial locality.



Memory Hierarchy

- The advantage of the locality can be achieved by implementing a memory hierarchy.
- Because of the differences in cost and access time, building a memory as a hierarchy is advantageous.





ENGINEERING THE FUTURE

ITU VLSI LABS

	Speed	CPU	Size	Cost (\$/bit)	Current Technology
CACHE	Fastest	Memory	Smallest	Highest	SRAM
MAIN MEMORY & RAM		Memory			DRAM
HARD DRIVE & DISK	Slowest	Memory	Biggest	Lowest	Magnetic Disk



Memory Hierarchy

- The idea is achieving the access time of the **fastest** memory component and yet having a memory size as large as the **largest** memory component.
- A large portion of the instruction and operand fetches are expected to be from the **cache** because it is fastest one.
- Some other portion of fetches not satisfied by the **cache** is performed by the **main memory**.
- The **hard drive** is accessed in the very infrequent cases in which a CPU instruction or operand fetch is not found in **main memory**.



Example

- **CPU :**
 - Memory address width 32 bits
 - Access time 1-ns clock cycle.
- **Main Memory:**
 - $2^{32} = 4$ GB main memory
 - 2 gb for instruction memory, 2 gb for data memory.
 - Access time 10-ns clock cycles.
- One-tenth speed is too slow.
- We can use a more efficient way.



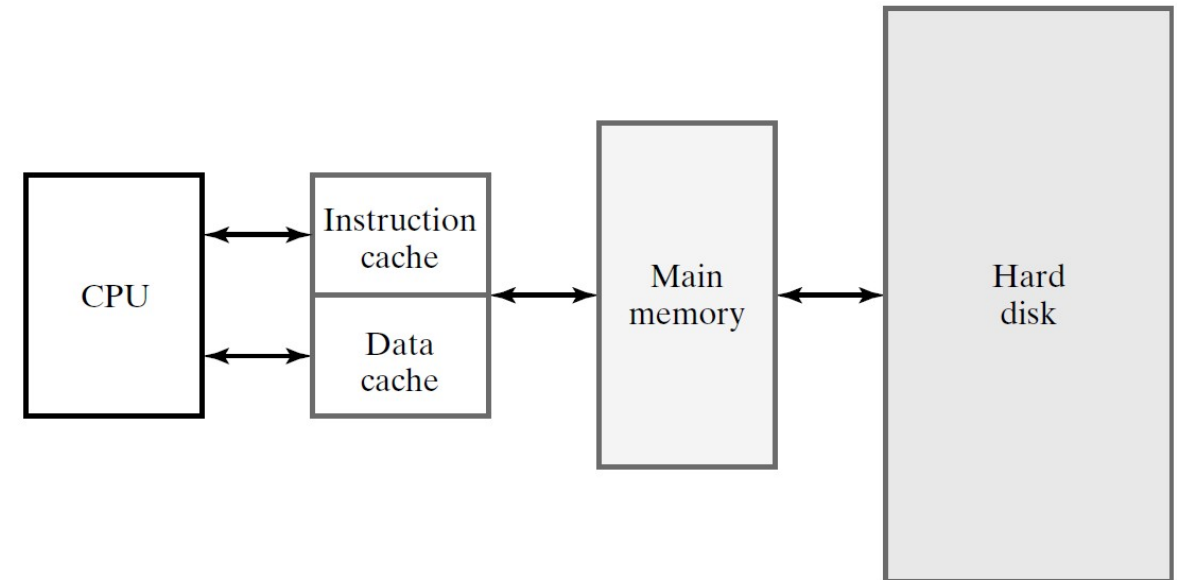
Example

- **Cache:**
 - 32 kb memory.
 - 16 kb for data-memory, 16 kb for instruction memory.
 - Fetch&Store 2-ns clock cycles.
- Suppose that %95 of the memory accesses performed from caches and remaining part %5 performed from main memory.
- Then the average access time is:
 - $\%95 \times 2 + \%5 \times 10 = 2.4 \text{ ns}$



Example

- This means that, on 19 out of every 20 the CPU operates at full speed, while the CPU will have to wait for 10 clock cycles for 1 out of every 20.
- This wait can be accomplished by stalling the CPU pipeline.





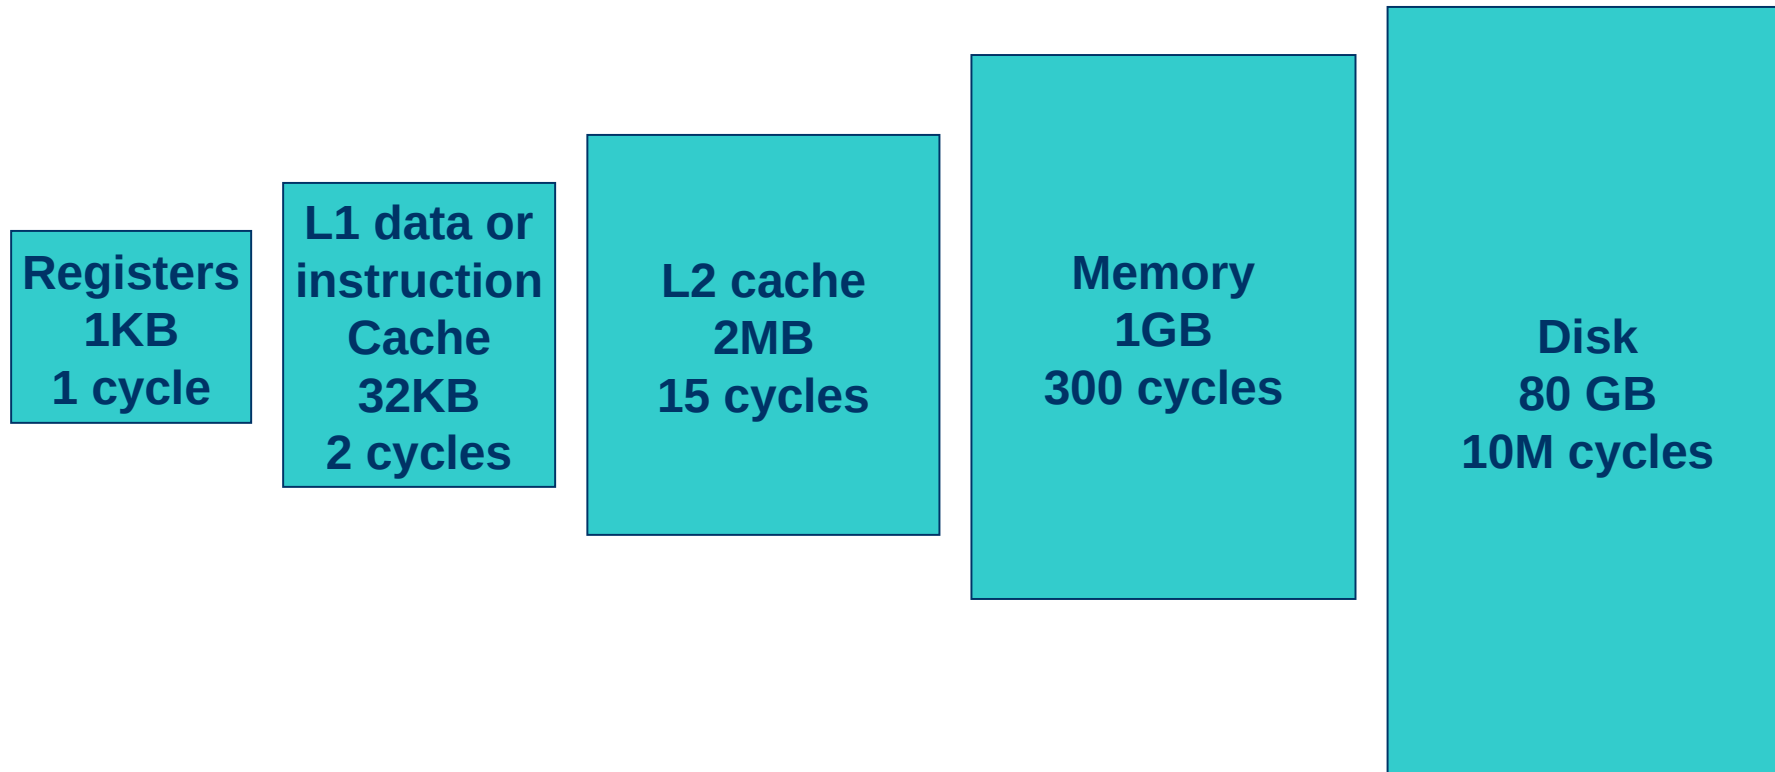
Example

- What about the size?
- Hard Disk:
 - 64 gb memory
 - Access time : 13ms = 1.3×10^7 ns
- Let's say %95 for cache %4.999995 for main memory.
- Then the average access time is:
 - $0.95 \times 2 + 0.04999995 \times 10 + 5 \times 10^{-8} \times 1.3 \times 10^7 = 3.05$ ns



Memory Hierarchy

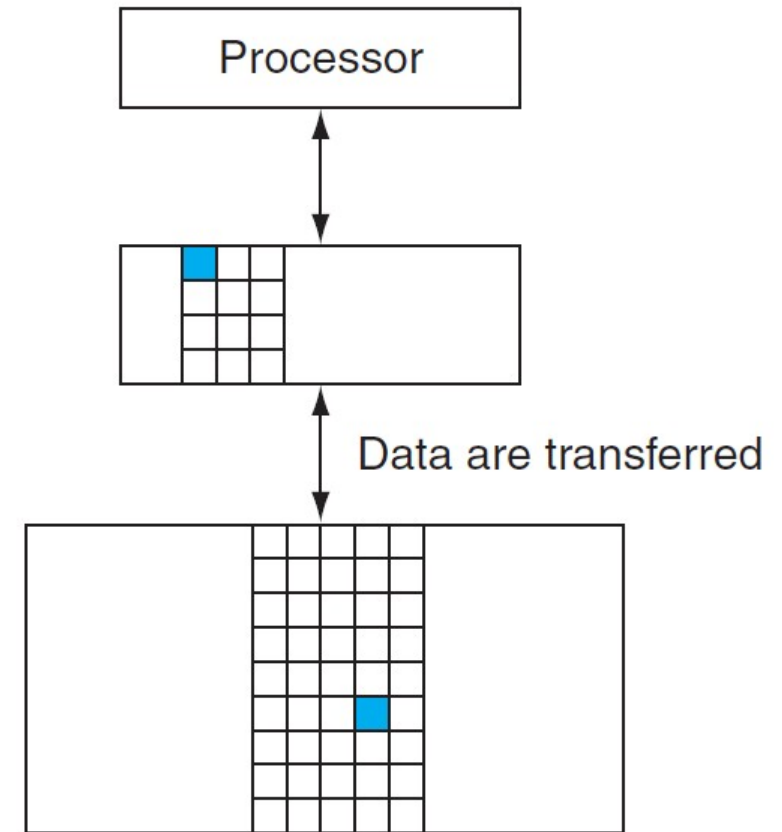
- As you go further, size and latency increase!





Cache

- A memory hierarchy can consist of multiple levels, but data is copied between **only** two adjacent levels at a time. This is called **two-level hierarchy**.
- The minimum unit of data that can be either present or not in a level is called a **block**.





Cache

- **Hit:** If the data requested by the processor appears in the cache (or some other level).
- **Miss:** If the data is not found in the cache, it is called a miss.
- **Hit Rate:** It is the fraction of memory accesses found in the cache, it is often used as a measure of the performance.
- **Miss Rate:** $(1 - \text{hit rate})$ is the fraction of memory accesses not found in the cache.



Cache

- **Hit Time:** The time needed to determine whether the access is a hit or a miss.
- **Miss penalty:** the time to replace a block in the cache with the corresponding block from the main memory.



Performance Metrics

- **Cache Capacity (C):** Memory size of the cache.
- **Block Size (b):** Number of data bytes can be transferred to the cache at one time.
- **Number of Blocks (B):** Number of blocks in the cache. ($B = C / b$)
- **Degree of Associativity(N):** Number of blocks in a set.
- **Number of Sets(S):** Number of words in a block. ($S = B / N$)



Cache

- Suppose that CPU requests an data item X_n .
- If X_n is initially in the cache. It is called **cache hit**.
- If not, this request results in a **miss**, and the word X_n is brought from memory into cache.

X_4
X_1
X_{n-2}
X_{n-1}
X_2
X_3

a. Before the reference to X_n

X_4
X_1
X_{n-2}
X_{n-1}
X_2
X_n
X_3

b. After the reference to X_n



Cache

- How do we know if a data item is in the cache?
- If it is, how do we find it?
- There must be information in the cache to identify the **address** of the word in **main memory!**

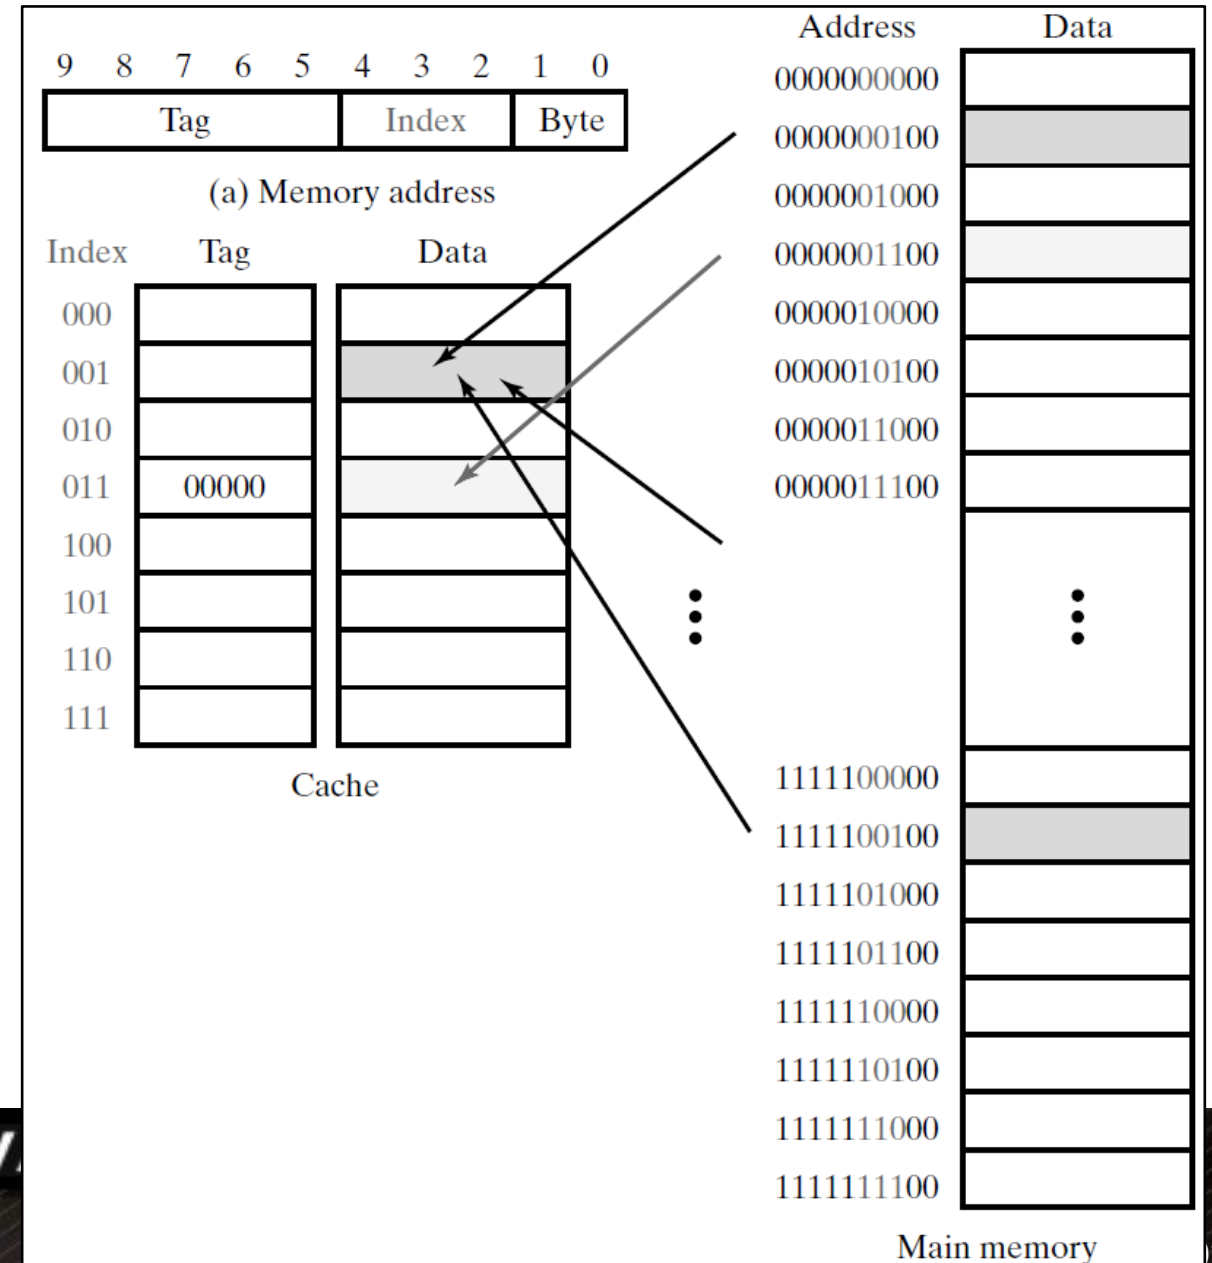


Direct Mapped Cache

- The simplest way to assign a location in the cache for each word in main memory is **direct mapping**.
- In a **direct mapped cache**, each main memory location is mapped directly to the cache.
- This cache structure is called **direct mapped cache**.

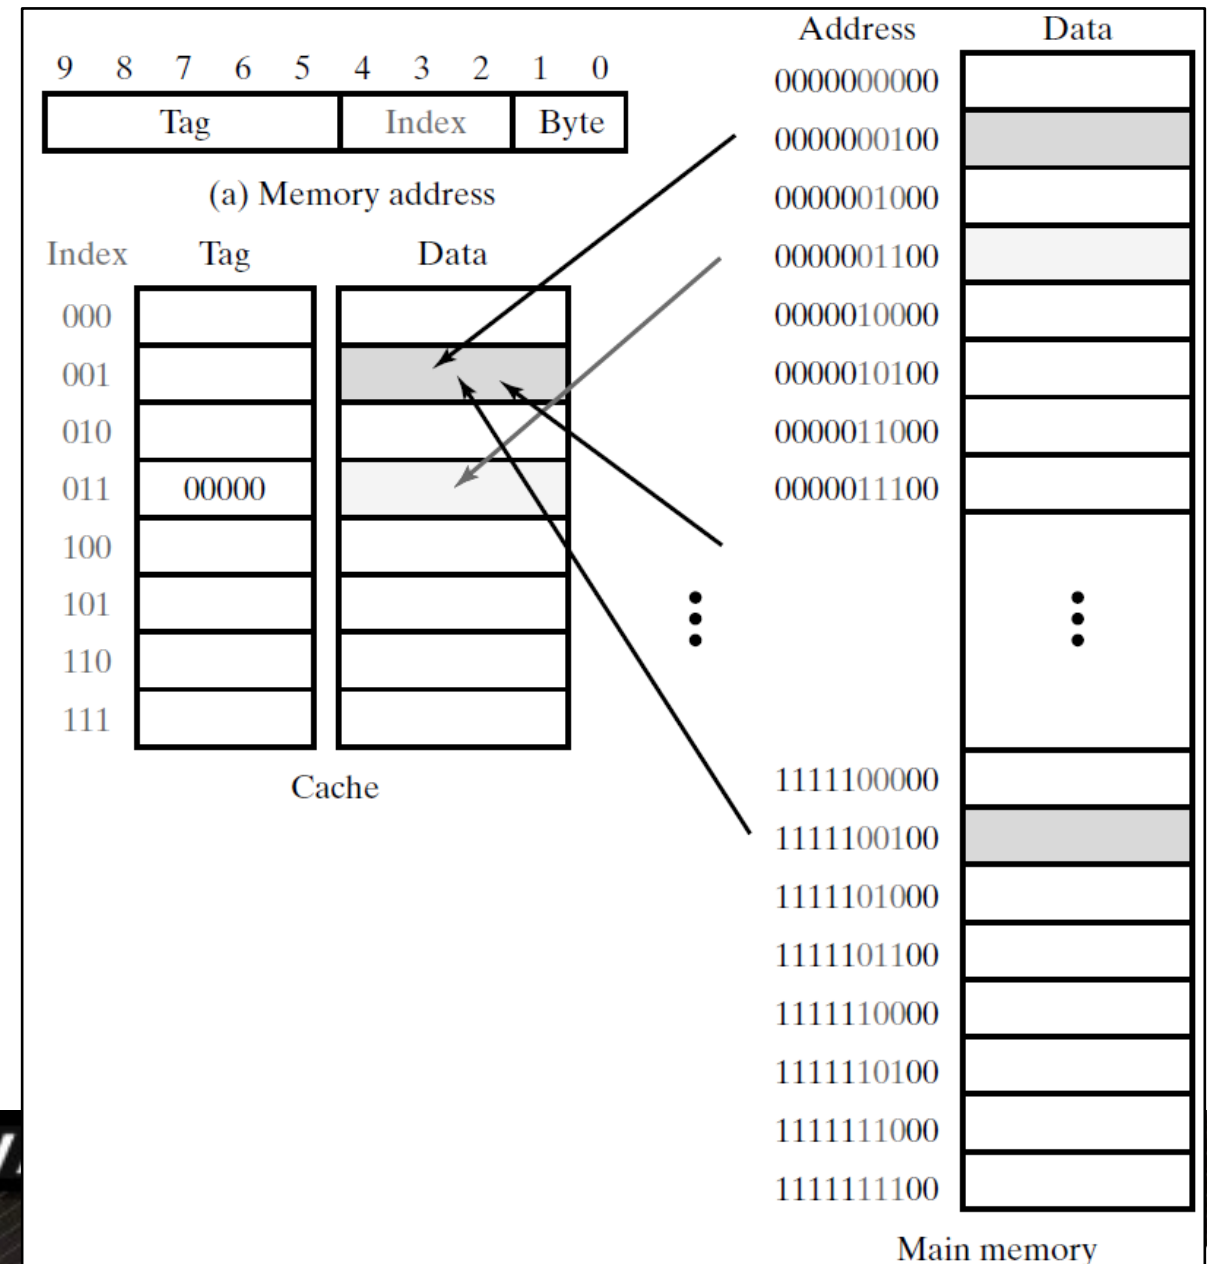


- **Memory size:** 256-words length (1kB)
- **Memory address:** 10-bit
- **Cache size:** 8-words length cache
- **Block size:** One block is a 32-bit word in the cache
- **Cache address:** 3-bit
- 2-4 of the main memory address is the **cache address**. It is called **index**.
- Upper 5 bits of the main memory address is the **tag**.





- The first two bits of the memory address is excluded because memory is **word-addressable** (block size = 32 bits)
- **Cache Index** = 3 bits, since the cache has 8-bit words length. $2^3 = 8$
- **Cache index** is used to refer to one block.
- **Tag** = 5 bits, this is the remaining part of the memory address.
- **Tag** is necessary to identify the word address in main memory.





Direct Mapped Cache

- Suppose that the CPU fetch an instruction from the location 000001100 in main memory.
 - This instruction may fetch from either the cache or main memory!
 - The cache separates the **tag** 00000 from the **cache address** 011.
 - If the **tag** is 00000, then the tags match, and the stored word fetched from cache.
- If the **tag** fetched from cache memory is not 00000, then there is a **tag mismatch**.
 - The cache control notifies main memory that it must provide the memory word.
 - In expectation of future accesses to the same memory address the **tag** 00000 and the word from memory is written in cache location 011.



Fully Associative Cache

- Direct mapping for a cache is not practical.
- Suppose that:
 - Instructions and data are in the same cache
 - There is a loop instruction in location 0000001100.
 - Besides, data from location 1111101100 is frequently used.
- **Instruction : tag = 11111, index = 011**
- **Data : tag = 00000, index = 011**




Fully Associative Cache

- In that case:
 - When the **instruction** in 0000001100 is fetched, location 011 in the probably contain the **data** from 1111101100 and tag 11111.
 - Then a cache miss occurs and tag 11111 would be replaced with tag 00000 in the cache.
 - But the next time the **data** is needed, another cache miss occurs, and the location in the cache would be replaced by the **instruction**.
 - Therefore, throughout the execution of the loop, both **instruction** fetch and **data** fetch cause many cache misses, significantly slowing CPU processing.

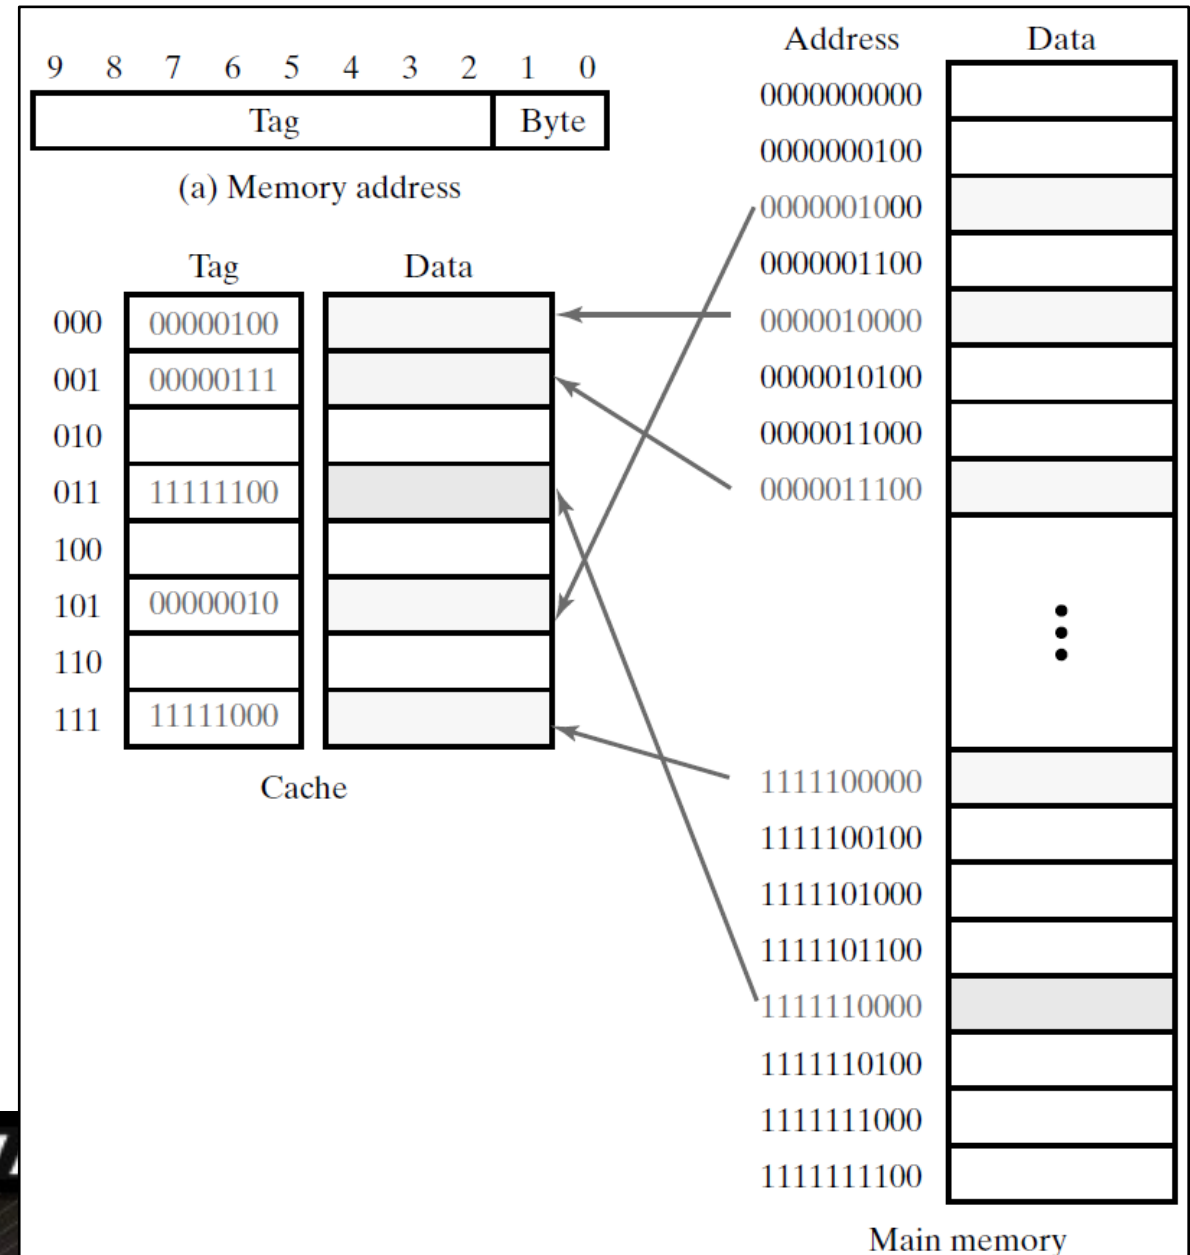


- In the **direct mapped cache**, only one of the 2^5 tag addresses can be mapped to the cache:
00000, 00001, 00010....11111
- In contrast, we can let locations in main memory map into an **arbitrary** location in the cache.
- It means any location in memory can be mapped to any one of the eight addresses in the cache.



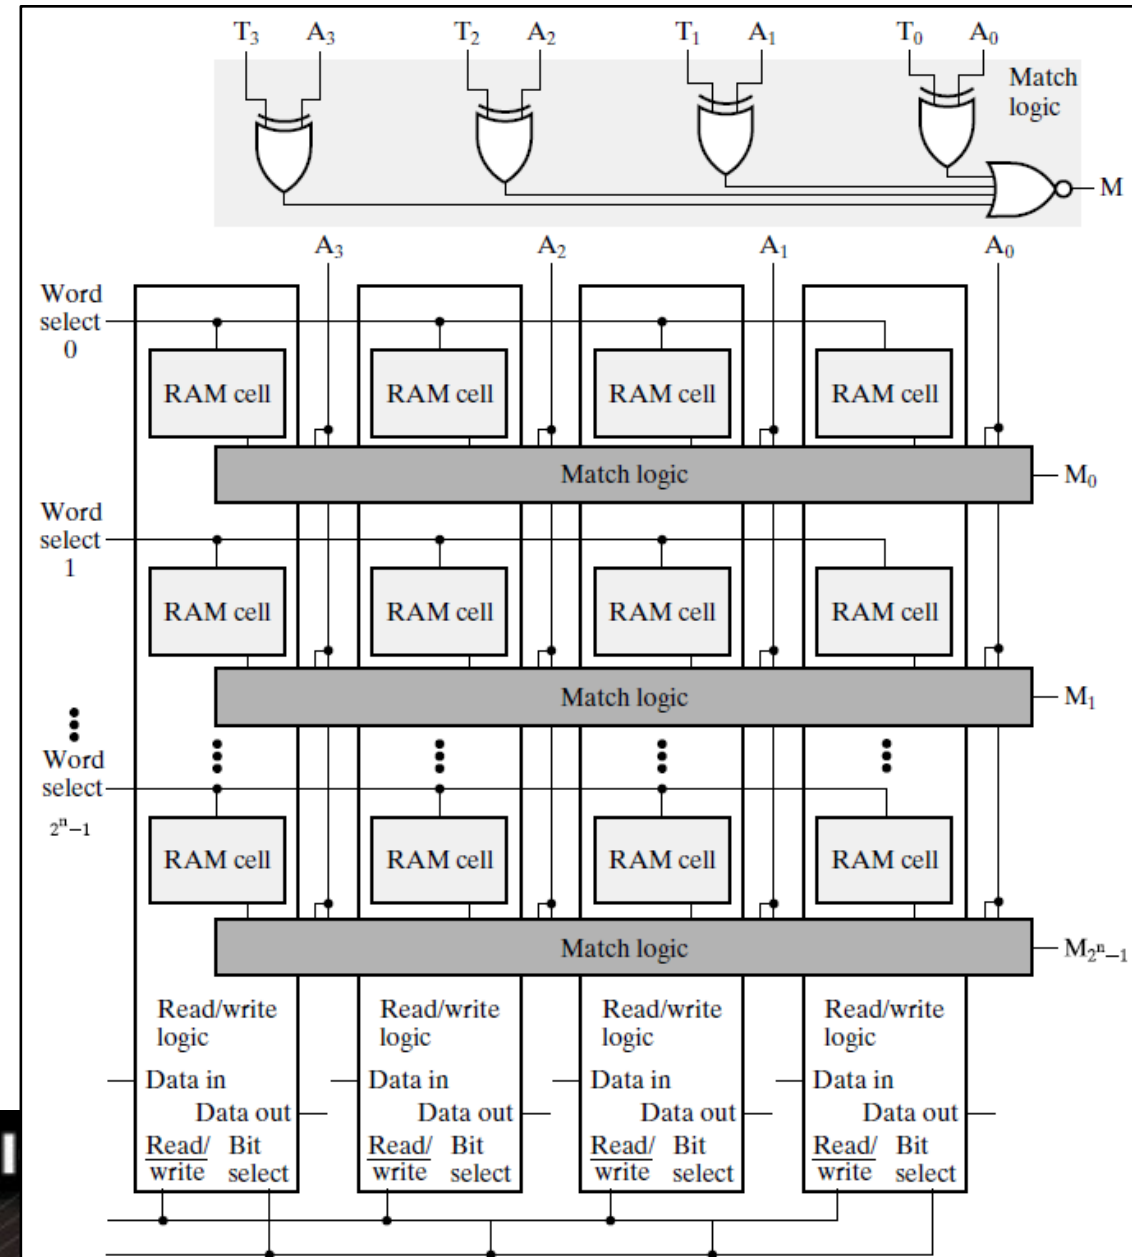


- **Tag** is now the full main memory word address.
- This is called **full associative mapping**.
- In a **fetch** process, **tags** are compared and if any match occurs then there is a cache hit, otherwise a cache miss.





- In **fully associative cache**, reading and comparing of tags can take too much time.
- Suppose that a fetch operation requires 2 ns, in the worst case 8 fetch is necessary and it takes 16 ns.
- Instead of **successive** reading of **tags**, associative memory structure (slicing model) can be used.
- Each tag storage row has match logic.





Fully Associative Cache

- The match logic does an equality comparison between the **tag T** and the applied address **A** from the CPU.
- Since all tags are unique, only two situations can arise in the associative memory:
 - there will be a match, with a 1 on the output of the match logic for one tag and a 0 on the remaining match logic outputs,
 - or there will be no match, and all of the match logic outputs will be 0.



Fully Associative Cache

- In a **cache miss** situation, the fully associative cache must bring the word and its address tag from the main memory for future accesses.
- Therefore, a replacement approach must be selected:
 - Random
 - FIFO : replace with the oldest entry
 - LRU : replace with the least used entry
- LRU approach is better than others.

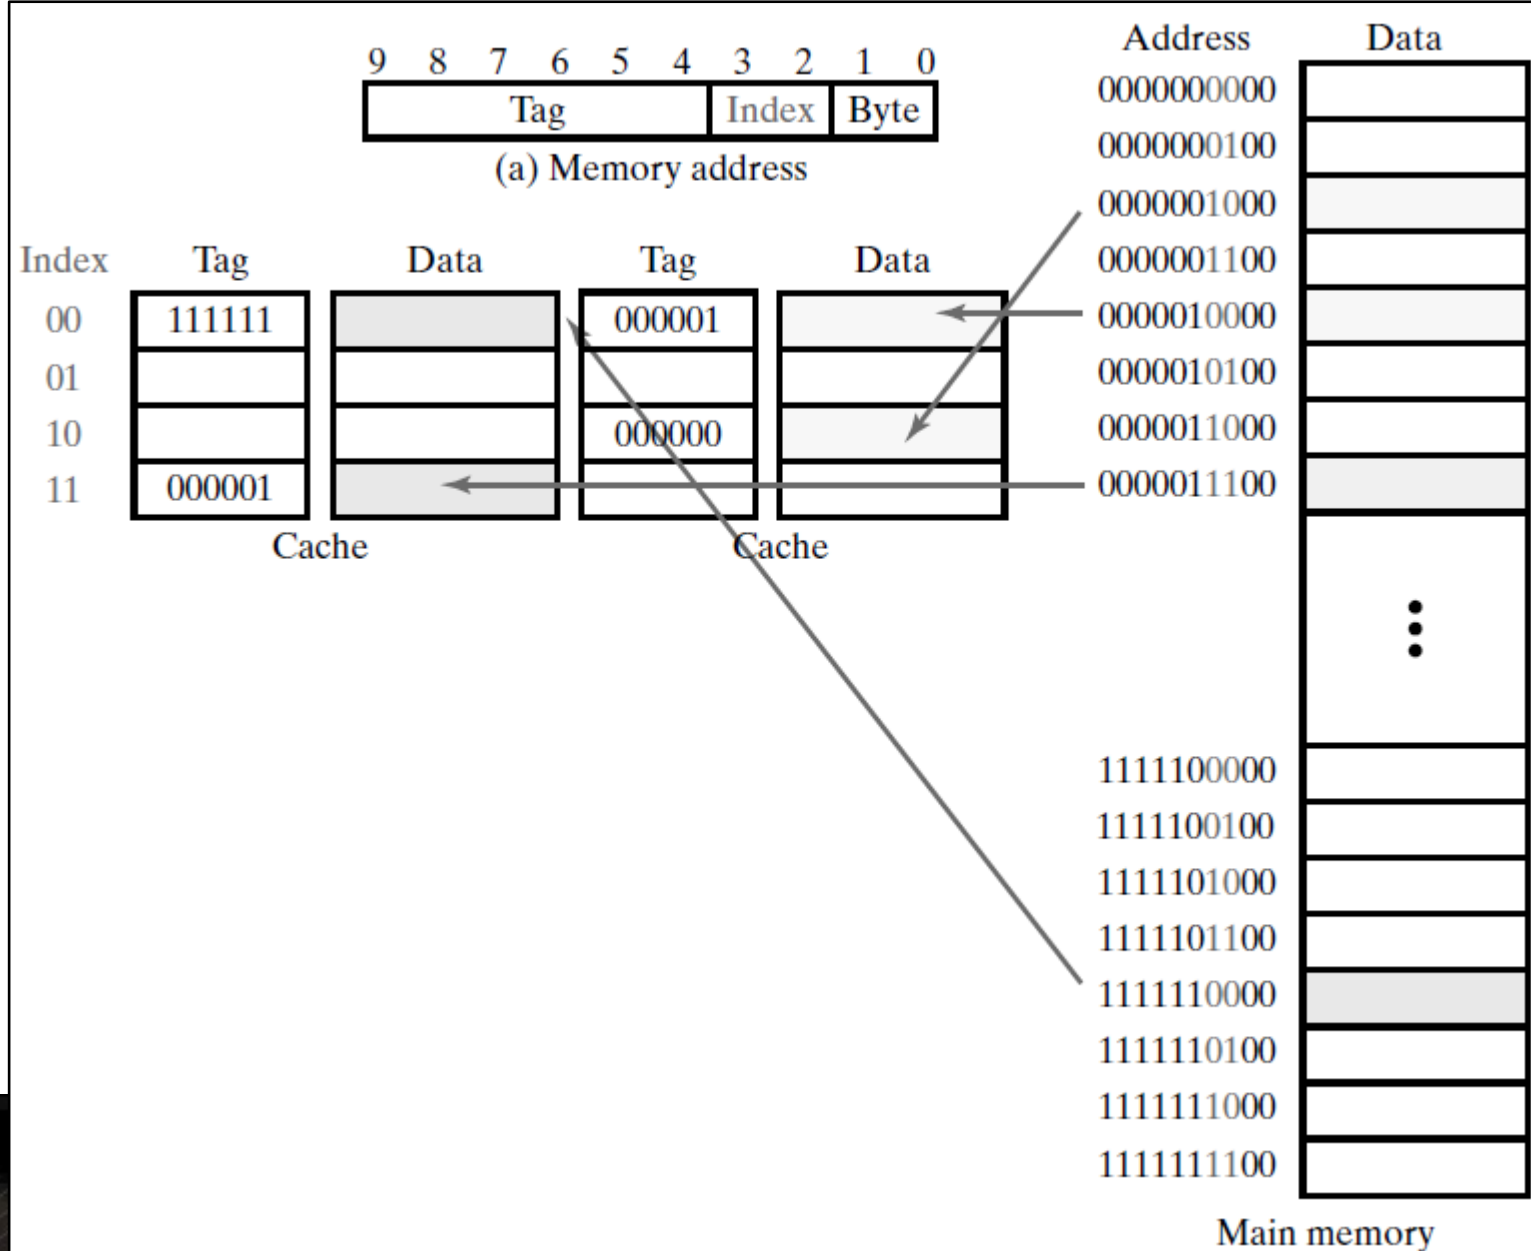


Set-Associative Mapping

- **Set-Associative Mapping** has better performance and eliminates the cost of most of the matching logic blocks.
- **Tags** and **words** from the cache are addressed by the lower-order address bits similar to **direct mapping** (index).
- Instead of having one location, there is a **set** of locations. If the set size equals **2**, then two tags and the two data words are read simultaneously.
- Tags are compared using just **2** math logic blocks.
- This is called **2-Way Associative Cache**.



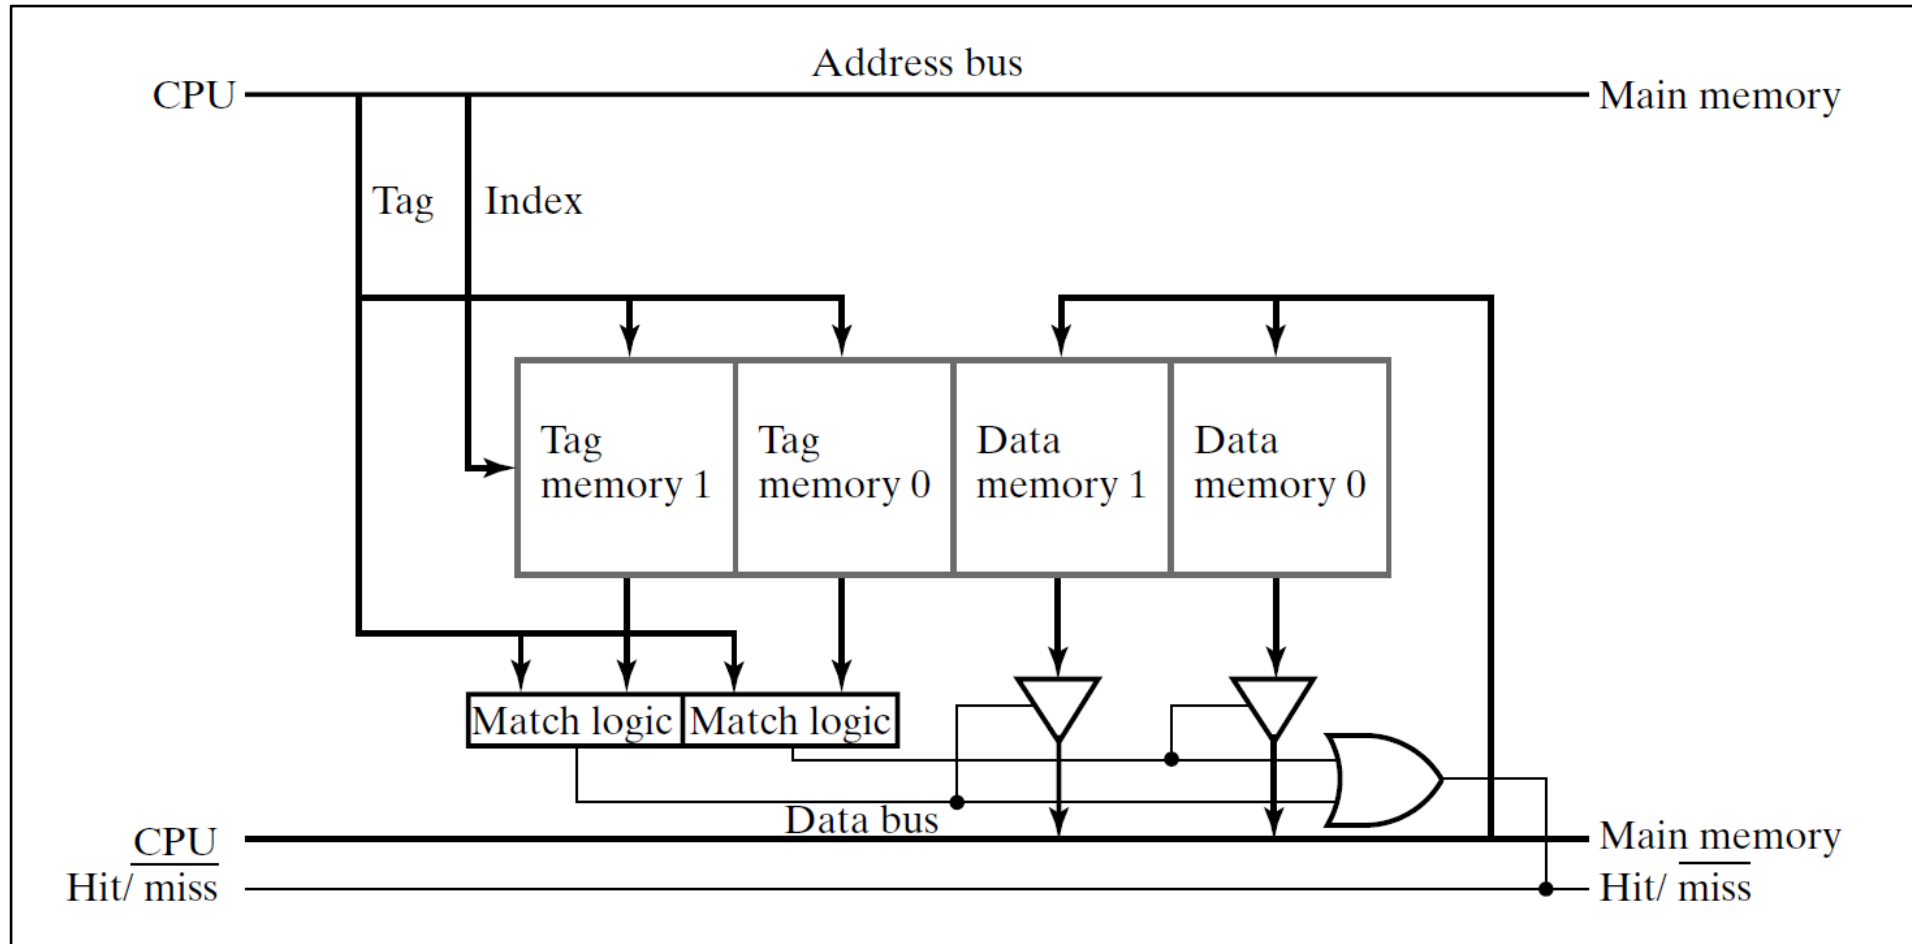
- Eight cache locations are arranged in four rows of two locations each.
- A cache entry can be in either of the two columns.
- **Index:** 2-bit
- **Tag:** 6 bits





2-Way Associative Cache

- 00 index for both cells are occupied by addresses 0000010000 and 1111100000.
- In order to accommodate another 00 index address like 1111100000, the **set size** would need to be at least three.
- Therefore, full associative cache is more **flexible** but 2-way associative cache has a **less complicated** hardware.





2-Way Associative Cache

- The index is used to address each row of the cache memory.
- The two tags read from the tag memories are compared with the tag to be fetch.
- If a match occurs, then the three-state buffer on the corresponding data memory output is activated, placing the data onto the data bus to the CPU.
- In addition, the match signal causes the output of the Hit/Miss.

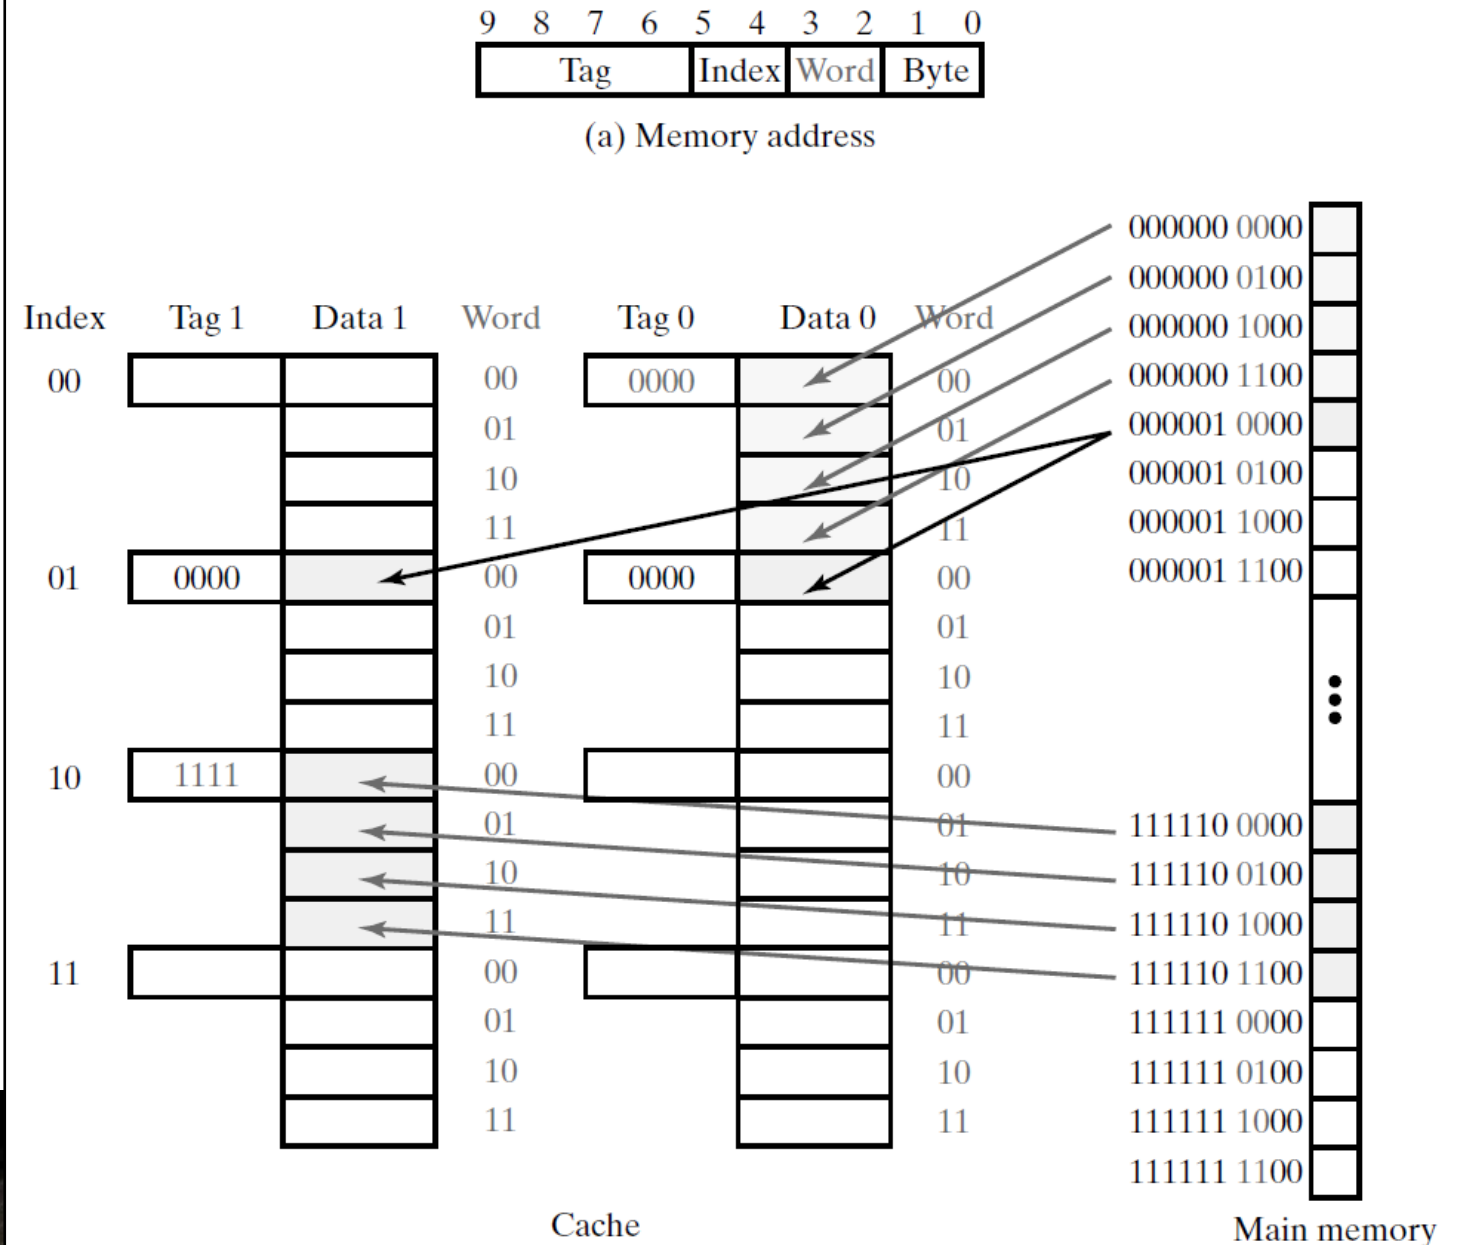


Line Size

- In real caches additional words close to the addressed one are also included in the cache.
- Therefore, when a cache miss occurs, a block of words called a **line** is fetched.
- By this way, spatial locality can be exploited.
- For example,
 - If four words are included in a line, then the addresses of the words in the line differ only in bits 2 and 3.
 - The use of a block of words changes the address fields.
 - The new field structure is **Word Field** including bits 2 and 3.



- **2-Way Associative Cache with 4-Word Lines**
- **Word** refers to one of the word in the set.
- **Index** refers to set.
- **Word:** 2 bits (4 words)
- **Index:** 2 bits (4 sets)
- **Tag:** 4 bits





2-Way Associative Cache with 4-Word Lines

- **Index** is used to read two **tags**.
- **Index** and the **Word address** are applied to read out two words from the two **tags**.
- 2 matching logic is used for comparison of the tags.
- If a match occurs, then the data word already read is placed on the memory bus to the CPU.
- Otherwise, a cache miss is signaled.



Path Size

- To load the entire line of words, the width of the memory bus is made more than one word wide.
- The path can be $4 \times 32 = 128$ bits wide. This allows the entire line to be placed in the cache in a single read cycle.
- If the path is narrower, then a sequence read cycles are required.
- A wide path can cause higher cost, but a narrower path can cause slow transfer of the line.



Cache Loading

- Before any words and tags have been loaded into the cache, all locations contain invalid information.
- As lines are fetched from main memory into the cache, cache entries become valid.
- To distinguish valid entries from invalid entries, a **valid bit** is used.
- A **valid bit** indicates whether an entry contains an valid address.



Cache Loading

Index	V	Tag	Data
000	N		
001	N		
010	N		
011	N		
100	N		
101	N		
110	N		
111	N		

a. The initial state of the cache after power-on

Index	V	Tag	Data
000	Y	10_{two}	Memory (10000_{two})
001	N		
010	Y	11_{two}	Memory (11010_{two})
011	N		
100	N		
101	N		
110	Y	10_{two}	Memory (10110_{two})
111	N		

d. After handling a miss of address (10000_{two})



Write Methods

- Write operations are:
 1. Write the result into main memory.
 2. Write the result into the cache.
 3. Write the result into both main memory and the cache.
- We have two categories:
 - write-through
 - write-back.



Write Methods

- In **write-through**, the result is always written to main memory.
 - This uses the main memory write time and can slow down processing.
 - The slowdown can be partially avoided by using **write buffering**.
 - In **write buffering**, address and word are stored in special registers called write buffers.
 - By this method, CPU can continue processing during the write to main memory.



Write Methods

- In **write-back method**, the result is written to the cache in case of a cache hit.
 - The goal of a write-back cache is to be able to write at the writing speed of the cache whenever there is a cache hit.
 - When a cache miss appears, the result is released from the cache and written into the main memory.
 - This avoids having all writes performed at the slower writing speed of main memory.



References

- D. A. Patterson, *Computer organization and Design*. San Francisco: Elsevier Science & Technology.
- M. M. Mano, and C. R. Kime,, *Logic and computer design fundamentals*. Boston: Pearson.